

Not All Instructions Are Forgotten Equal

Tomás P. Korenblit 

Universidad Nacional de San Martín, Buenos Aires, Argentina
tomaskorenblit@gmail.com

Abstract. LLM agents increasingly run for hundreds of turns under instructions given once at the start, safety guardrails among them, that must hold as the context keeps growing. Yet we find these instructions are not retained equally: a model keeps following some and drops others, with nothing to flag the lapse. Prior work reports mean compliance and reinforces every instruction uniformly; we measure retention per instruction instead. We plant preferences in multi-turn coding sessions on three open-source Python codebases, across five models and twelve instruction types, and grade compliance with deterministic checkers rather than an LLM judge, so the scores are reproducible. A Bayesian hierarchical model then estimates how strongly each instruction is retained. Most preferences are already followed by default, and a few only when explicitly restated. Only three of twelve types gain enough from restating to justify it (two under a single-model robustness check), so blanket reinforcement spends most of its tokens for no gain. The unevenness cuts the other way for safety: a guardrail is just another instruction, and the rules a model would not follow on its own are exactly the ones a single statement may fail to secure. Since safety is conjunctive, an aggregate compliance rate cannot tell whether the rule that matters still holds, and in an agent loop nothing flags the lapse without per-instruction monitoring.

Keywords: instruction retention , context rot , Bayesian inference , AI safety , LLM evaluation

No Todas las Instrucciones se Olvidan Igual

Abstract. Los agentes LLM operan cada vez más durante cientos de turnos bajo instrucciones dadas una sola vez al inicio, las barreras de seguridad entre ellas, que deben sostenerse a medida que el contexto crece. Sin embargo, encontramos que esas instrucciones no se retienen por igual: el modelo sigue cumpliendo unas y abandona otras, sin nada que señale la falla. El trabajo previo reporta cumplimiento promedio y refuerza cada instrucción de manera uniforme; nosotros medimos la retención por instrucción. Plantamos preferencias en sesiones de programación multi-turno sobre tres repositorios de Python de código abierto, en cinco modelos y doce tipos de instrucción, y evaluamos el cumplimiento con verificadores determinísticos en lugar de un juez LLM, de modo que los

puntajes son reproducibles. Un modelo jerárquico bayesiano estima cuánto se retiene cada instrucción. La mayoría de las preferencias ya se cumplen por defecto, y unas pocas solo al enunciarse de forma explícita. Solo tres de doce tipos ganan lo suficiente al reenunciarse para justificarlo (dos bajo una verificación con un solo modelo), de modo que el refuerzo uniforme gasta la mayoría de sus tokens sin ganancia. La disparidad juega en contra de la seguridad: una barrera de seguridad es una instrucción más, y las reglas que el modelo no seguiría por sí solo son justamente las que un solo enunciado puede no asegurar. Como la seguridad es conjuntiva, una tasa de cumplimiento agregada no puede decir si la regla que importa sigue vigente, y en un loop de agente nada advierte la falla sin monitoreo por instrucción.

Keywords: retención de instrucciones , degradación de contexto , inferencia bayesiana , seguridad de la IA , evaluación de LLMs

1 Introduction

Language models are increasingly deployed as autonomous agents that run in loops, call tools, and act over hundreds of turns with no human reviewing each step. An agent is governed by its instructions: the operational limits and safety guardrails it is given, usually once and at the start, that must hold for the whole run. Yet adherence to instructions erodes as the context grows, a phenomenon we refer to as *context rot*. Laban et al. (2025) measured average compliance drops of 39% across 15 models in multi-turn settings. The cause is architectural: attention is not uniformly distributed across the context window, and information in middle positions receives less weight than at the boundaries (Chowdhury, 2026; Liu et al., 2024). An instruction stated once and then buried under accumulating tool output competes for attention against everything that follows it, and in a deployed loop there is no user to notice the lapse and restate it.

This is a safety concern before it is a cost concern. When the model drops a preference, a chat user notices and restates it; a deployed agent has no such supervisor, so a guardrail that decays does so silently while the agent keeps acting through tools whose effects are real and often irreversible (Mu et al., 2025). We plant ordinary coding preferences rather than guardrails, but the two are the same kind of object (an instruction stated once and then left to compete for attention), so how well a preference survives is a measurable stand-in for how a stated guardrail would hold. Repeating every instruction each turn, the standard mitigation, multiplies prompt cost and still cannot tell an operator which constraint remains in force. Because safety is conjunctive, an aggregate compliance rate says nothing about whether the one rule that matters is among those that lapsed.

No prior work has measured this heterogeneity at the level of individual instructions. Existing studies report mean compliance across instruction sets

(He et al., 2024; Laban et al., 2025) or propose uniform mitigation strategies such as periodic reminders (Dongre et al., 2025). But if instructions are retained unequally, the mean hides the per-instruction differences a reinforcement policy would act on. An instruction that is already retained does not need reinforcement, and one that is harmed by repetition should not receive it; without per-instruction estimates, a policy cannot tell these cases apart.

Bayesian Knowledge Tracing (Corbett & Anderson, 1994) offers a framework for exactly this kind of estimation. Originally developed in education, BKT estimates the probability that a student has mastered a concept based on a sequence of observed responses. We adapt this framing: the LLM is the student, and its instructions are the concepts. A BKT-based reinforcement system would only be useful if instruction retention is in fact heterogeneous. We test this using a Bayesian ordered logistic model with hierarchical effects per decision type, fitted to compliance scores collected from realistic multi-turn coding conversations across three open-source codebases.

We report three findings. First, retention heterogeneity across instruction types is large, with treatment effects on the log-odds scale ranging from -2.4 to $+5.4$ (group-level standard deviation $\sigma_\beta = 2.1$). Second, one instruction type (`dependencies_no_new`) shows a marginal negative effect (its 94% HDI barely excludes zero), which we read as a likely measurement artifact rather than a robust harm. Third, within the resolution of this design, the model and codebase size explain little variance in retention; because baseline data comes from a single model, we read this as model-invariance within our design rather than an established absence of model effects. A Bayesian model can then identify which instructions to reinforce and which to leave alone.

2 Related work

Instruction compliance degradation. Several recent studies have established that LLMs lose adherence to instructions as conversations grow. Laban et al. (2025) documented an average 39% compliance drop across 15 models in multi-turn settings over 200K conversations. He et al. (2024) showed that accuracy on multi-instruction tasks drops from 87.7% to 70.7% within three turns, even when instructions are given simultaneously. Du et al. (2025) demonstrated that increasing the context window alone does not resolve this degradation. All three studies report aggregate metrics across instruction types. None examines whether different instructions are retained to different degrees.

Positional attention and prompt saturation. The mechanism behind compliance degradation has roots in the attention architecture itself. Liu et al. (2024) identified a U-shaped attention curve: information at the beginning and end of the context receives more weight, while middle positions are underweighted. This positional bias explains why system prompt instructions, which occupy position zero, remain more stable than user-stated preferences buried mid-conversation. Mu et al. (2025) showed a related saturation effect: as the number of guardrails

in a system prompt increases, compliance with any individual rule drops sharply. Adherence therefore degrades unevenly: it depends on where an instruction sits and how many others surround it.

Uniform mitigation strategies. Existing approaches to maintaining compliance treat all instructions equally. Dongre et al. (2025) modeled instruction drift as a bounded stochastic process and proposed periodic reminders to reduce divergence. Leviathan et al. (2025) studied the effect of prompt repetition on instruction stability. Both approaches reinforce all instructions at each intervention, regardless of whether an individual instruction has been retained or forgotten. This uniform strategy becomes expensive at scale and, as our results show, can be counterproductive for instructions where repetition degrades compliance.

Bayesian Knowledge Tracing. BKT has tracked per-concept mastery probabilities in education for over 30 years (Corbett & Anderson, 1994). The per-instruction heterogeneity that makes BKT useful there also appears in LLM instruction compliance, with each instruction in the role of a concept.

3 Method

3.1 Decision planting protocol

We simulate an AI coding assistant working on a real open-source codebase over 25 turns. At each turn, the assistant receives a user message and may call tools (read files, write code, search the codebase). Real source files are injected into the conversation at turns 0–19, building up context pressure as the conversation progresses.

In the **treatment** condition, 12 coding preferences are embedded as casual inline requests within the user’s messages at specific turns. Six are planted early (turns 1–7) and six late (turns 14–19). Each preference targets a different coding practice: code style (NumPy broadcasting, list comprehensions), architecture (subclassing vs. standalone functions), testing (parametrize, approximate assertions), naming (underscore prefix, no abbreviations), dependencies (no new imports, module-level constants), and documentation (NumPy-style docstrings, regex comments). The instructions are phrased as natural user preferences, not system prompt rules (e.g., “Important: when writing array operations for this project, always use NumPy broadcasting instead of for loops”).

In the **baseline** condition, the conversation follows the same structure with the same file injections and the same task, but no preferences are planted. This allows us to measure what the model does by default for each decision type.

At turns 20–25, the user requests code generation tasks designed to elicit each decision. Each test turn evaluates exactly two decisions (one early-planted, one late-planted), and the model’s code output is scored by deterministic checkers.

3.2 Compliance measurement

Each of the 12 decision types is evaluated by a deterministic checker that parses the model’s code output using regular expressions and pattern matching. Checkers are fully automated with no LLM-in-the-loop evaluation. This makes results reproducible across runs. Each checker produces an ordinal score on a 0–3 scale:

- **0**: Completely ignored or violated
- **1**: Minimal or inconsistent compliance
- **2**: Mostly followed with minor deviations
- **3**: Fully followed

For example, the `parametrize` checker searches for `pytest.mark.parametrize` decorators and penalizes separate test functions. The `broadcasting` checker detects `for` loops and rewards NumPy operations. The `no_new_imports` checker counts `import` statements in generated code.

3.3 Bayesian model

We model the ordinal compliance scores using an ordered logistic regression (cumulative link model) with hierarchical effects per decision type. The ordered logistic is appropriate because the 0–3 scores have unequal intervals: the gap between 0 (ignored) and 1 (minimal compliance) is qualitatively different from the gap between 2 (mostly followed) and 3 (fully followed). A linear model would incorrectly assume equal spacing.

For each observation i , the latent compliance η_i is:

$$\eta_i = \alpha_{d[i]} + \beta_{d[i]} \cdot \text{treatment}_i \quad (1)$$

where $d[i]$ indexes the decision type, α_d is the baseline intercept (compliance without being told), and β_d is the treatment effect (how much telling improves compliance). The observed score is determined by where η falls relative to three ordered cutpoints $c_1 < c_2 < c_3$.

Both α and β are given hierarchical priors with non-centered parameterization:

$$\alpha_d = \mu_\alpha + \sigma_\alpha \cdot \tilde{\alpha}_d, \quad \tilde{\alpha}_d \sim \text{Normal}(0, 1) \quad (2)$$

$$\beta_d = \mu_\beta + \sigma_\beta \cdot \tilde{\beta}_d, \quad \tilde{\beta}_d \sim \text{Normal}(0, 1) \quad (3)$$

The hyperpriors are weakly informative on the log-odds scale: $\mu_\alpha \sim \text{Normal}(0, 1.5)$, $\mu_\beta \sim \text{Normal}(0, 1)$ (centered at zero, no prior assumption about direction), $\sigma_\alpha, \sigma_\beta \sim \text{Exp}(1)$. Cutpoints use an offset parameterization: $c_1 \sim \text{Normal}(0, 1.5)$ with positive increments $\Delta c \sim \text{Exp}(1)$ to guarantee ordering.

The key scientific parameter is σ_β : the group-level standard deviation of treatment effects. A large σ_β indicates that different instructions respond differently to being stated, which is the premise required for selective reinforcement.

We initially fit a full model including per-model and per-codebase intercepts. Both group-level standard deviations were small, with posteriors concentrated near zero ($\sigma_{\text{model}}, \sigma_{\text{codebase}} \approx 0.3$, 94% HDI upper bound below 0.7), so these terms add no explanatory power. We therefore report the simpler per-decision model, whose diagnostics are clean (zero divergences, all $\hat{R} \leq 1.01$, minimum ESS = 1311). With baseline data from a single model (Section 6.3), per-model effects are in any case weakly identified.

The model was implemented in PyMC (Abril-Pla et al., 2023). Posterior samples were drawn using the No-U-Turn Sampler (NUTS; Hoffman & Gelman, 2014) via the nutpie sampler, a Rust-based NUTS implementation (Seyboldt, 2024), with 4 chains of 1000 draws each following 1000 tuning steps. Prior predictive checks confirmed that the priors produce plausible score distributions before fitting. Posterior predictive checks verified that the fitted model reproduces the observed data.

4 Experimental setup

4.1 Codebases

We selected three open-source Python libraries from the Bayesian statistics ecosystem, chosen to span a range of codebase sizes and thus context pressure levels:

- **Bambi** (~10K lines): a high-level interface for Bayesian regression. By turn 20, injected files produce approximately 98K tokens of context.
- **ArviZ** (~25K lines): a library for exploratory analysis of Bayesian models. Context reaches approximately 160K tokens by turn 20.
- **PyMC** (~57K lines): a probabilistic programming framework. Context reaches approximately 203K tokens by turn 20.

All repositories were cloned at fixed commits: Bambi (5110615, 2026-03-17), ArviZ (`arviz-base` at 374cd28 and `arviz-stats` at 8d6316a, 2026-03-26/29), and PyMC (`bca2b1e`, 2026-03-27, branch v6). Source files were injected into the conversation at turns 0–19 via the tool execution system, which serves real file contents from the cloned repositories.

4.2 Models

We evaluated five instruction-following models across two experiment phases. Table 1 summarizes the models and conditions.

Qwen served as the primary model with both baseline and treatment conditions across all three codebases (18 conversations). The remaining four models were tested in the treatment condition only, to verify that the retention patterns generalize across architectures. Gemma was tested on Bambi only due to budget constraints.

Table 1. Models and experimental conditions. Only Qwen was tested in both baseline and treatment conditions.

Model	Parameters	Baseline	Treatment
Qwen 3.5	27B	✓	✓
Gemini 3.1 Flash Lite	–		✓
GPT-5.4-nano	–		✓
Nemotron 3 Super	120B		✓
Gemma 4	26B		✓

4.3 Dataset

We ran 28 conversations in total: 9 baseline and 19 treatment. Models were accessed through OpenRouter (Qwen, Gemini, Gemma, Nemotron) and the OpenAI API (GPT-5.4-nano), queried at each provider’s default sampling settings (temperature not overridden), one generation per turn capped at 2048 (Qwen) or 4096 tokens; scores reflect a single draw, not an average. Not every conversation produced a scorable response at every test turn; the final dataset consists of 244 compliance observations: each observation is one decision type scored at one test turn in one conversation. Observations span 5 models, 12 decision types, 3 codebases, and 2 conditions. The baseline condition contributes 70 observations (all from Qwen), and the treatment condition contributes 174 observations across all five models. All conversation logs, checker code, and analysis notebooks are available in the supplementary material.

5 Results

The model converged cleanly: zero divergences, $\hat{R} \leq 1.01$ for all parameters, and a minimum effective sample size of 1311. Posterior predictive checks confirm that the model reproduces the observed score distribution (Figure 1).

5.1 Treatment effect heterogeneity

The group-level standard deviation of treatment effects is $\sigma_\beta = 2.11$ (94% HDI: [1.06, 3.28]), indicating that different instruction types respond differently to being stated. Individual treatment effects on the log-odds scale range from -2.36 to $+5.44$ (Figure 2).

The posterior separates the 12 decision types into three groups:

Instructions that benefit from reinforcement. Three decision types have treatment effects with 94% HDI entirely above zero. `testing_parametrize` has the largest effect ($\beta = 5.44$, HDI: [2.94, 8.10]): models almost never use `pytest.mark.parametrize` unprompted (baseline mean: 0.00) but retain the preference when told (treatment mean: 2.79). `dependencies_module_constants` ($\beta = 2.76$, HDI: [0.60, 5.00]) and `architecture_standalone` ($\beta = 2.38$, HDI: [0.29, 4.63]) follow the same pattern: near-zero baseline, meaningful retention when stated.




Fig. 1. Posterior predictive check. Left: the model reproduces the observed score distribution across all four ordinal categories. Right: the model captures the difference in mean compliance between baseline (1.70) and treatment (2.16) conditions.

A marginal negative effect. The only decision type with a negative estimate is `dependencies_no_new` ($\beta = -2.36$, 94% HDI $[-4.46, -0.03]$). Its magnitude is uncertain, since the 94% HDI only just excludes zero, though the negative direction itself is credible ($P(\beta < 0) = 0.98$). The baseline mean score is 2.25 (models naturally avoid adding imports), while stating “do not add new imports” lowers the treatment mean to 0.77. One hypothesis is that the explicit negation confuses the model’s instruction following. But the checker for this type (`no_new_imports`) counts *any* import statement as new, including re-shown imports already present in the file (Section 6.3), so we treat this as a candidate measurement artifact rather than a robust harm. Isolating the negation from the content (comparing “avoid new imports” with “use only existing imports”) would settle it.

Instructions that need no intervention. The remaining eight decision types show treatment effects with HDIs crossing zero. Several of these have high baseline compliance (`broadcasting`: 3.00, `numpy_style`: 3.00, `snake_no_abbrev`: 2.67), meaning models already follow these conventions without prompting; reinforcing them does not change the score.

5.2 Model and codebase effects

As described in Section 3.3, a model with per-model and per-codebase intercepts placed both group-level standard deviations near zero, so we report the per-decision model. Treatment-condition compliance is consistent across the five models (26B to 120B for the three with disclosed sizes) and three levels of context pressure (98K to 203K tokens), and no per-model or per-codebase effect is detectable at this resolution. With a single-model baseline and shrinkage priors, though, this is weak evidence of model-invariance, not a demonstration, and per-model retention cannot be fully separated (Section 6.3). One reading we can set aside, though: retention does not climb from the 26B to the 120B model, so




Fig. 2. Per-decision treatment effects (β_d) with 94% HDI. Dark bars: 94% HDI excludes zero. Hatched: negative effect. Light bars: inconclusive.

the heterogeneity does not appear to be model capability under another name. A larger model is not automatically a more compliant one, which is what we would need to see if these differences were a capability effect rather than a property of the instruction.

5.3 Reinforcement priorities

We use the posterior to derive a reinforcement ranking. For each decision type d , we compute $P(\text{score} \geq 2 \mid \text{told})$ and $P(\text{score} \geq 2 \mid \text{not told})$ across all posterior samples, and define the reinforcement gain as their difference (Figure 3). This is not a closed-loop policy evaluation; it is a posterior-derived priority ranking that identifies which instructions would benefit from reinforcement.

Table 2 summarizes the ranking. Of 12 instruction types, only 3 have a reinforcement gain with 94% HDI entirely above zero. A uniform policy that reinforces all 12 would spend tokens on 8 instructions that need no help and 1 with a marginal negative effect. This ranking is a cost criterion, not a safety one: a low gain may mean the model already complies, or that the checker saw nothing (a default-2 ceiling), neither of which licenses dropping a safety-relevant rule. It saves tokens on ordinary preferences; it does not choose which guardrails to stop stating.




Fig. 3. Expected change in $P(\text{score} \geq 2)$ from reinforcing each decision, with 94% HDI. Dark bars: HDI excludes zero. Hatched: reinforcement is harmful.

6 Discussion

6.1 What determines retention

The treatment effects follow a pattern: instructions that ask for something the model would not do by default are retained when stated, while instructions that match existing behavior show no treatment effect. The three decisions with clear positive effects (**parametrize**, **module_constants**, **standalone**) all have baseline compliance near zero. The decisions with no treatment effect (**broadcasting**, **numpy_style**, **snake_no_abbrev**) have baseline compliance near 3.0. The models’ defaults already match these conventions, so restating them has no effect.

The negative effect on **dependencies_no_new** is the only one in the study. We call it marginal, since its 94% HDI barely excludes zero, though its direction is credible ($P(\beta < 0) = 0.98$); the magnitude, and whether it reflects negation framing or the checker artifact (Section 6.3), remain open. The instruction “do not add new imports” is the only one phrased as a negation. All other instructions tell the model what to do, not what to avoid. This raises the question of whether the negative treatment effect reflects something about the content of the instruction (import management) or its framing (negation). Which case holds determines the fix: if negation is the problem, reformulating as “use only existing imports” might produce a positive effect. If import management is inherently hard to retain, no framing will help. Distinguishing these hypotheses requires a controlled comparison that holds semantic content constant while varying the negation frame. We adopt this comparison in the future work below.

Table 2. Reinforcement priorities derived from posterior estimates. Gain is $\Delta P(\text{score} \geq 2)$ from reinforcement. The five skipped decisions are `broadcasting`, `listcomp`, `architecture_extend`, `snake_no_abbrev`, and `numpy_style`.

Decision type	Gain	94% HDI Action
<code>testing_parametrize</code>	+0.78 [+0.46, +0.96]	Reinforce
<code>module_constants</code>	+0.51 [+0.18, +0.76]	Reinforce
<code>arch_standalone</code>	+0.34 [+0.06, +0.61]	Reinforce
<code>naming_underscore</code>	+0.24 [-0.09, +0.53]	Uncertain
<code>testing_assert_almost</code>	+0.16 [-0.01, +0.40]	Uncertain
<code>docs_regex_comments</code>	+0.14 [-0.12, +0.45]	Uncertain
5 other types	< +0.04	crosses zero Skip
<code>deps_no_new</code>	-0.47 [-0.77, -0.05]	Likely artifact

6.2 Robustness checks

To address the concern that baseline data comes from Qwen only, we fit the same model restricted to Qwen observations (132 of 244, both conditions). The per-decision treatment effects correlate at $r = 0.80$ with the full-model estimates, and the group-level heterogeneity parameter remains stable ($\sigma_\beta = 2.35$ vs. 2.11 in the full model). Three of the strongest effects replicate: `parametrize` ($\beta = 6.13$), `module_constants` ($\beta = 2.25$), and `dependencies_no_new` ($\beta = -1.07$, still negative though with wider intervals).

One decision, `architecture_standalone`, does not replicate in the Qwen-only analysis ($\beta = 2.38$ in the full model vs. $\beta = -0.71$ with Qwen alone). Its treatment effect in the full model appears driven by the other four models. We reduce the count of clearly beneficial decisions from three to two when considering only Qwen data.

Leave-one-out cross-validation shows no problematic observations (all Pareto $\hat{k} < 0.7$, ELPD = -228.8 ± 14.5). Per-decision posterior predictive checks confirm that the model reproduces the observed mean score for all 12 decision types within the 90% predictive interval (Figure 4).

6.3 Limitations

Baseline data comes from one model only (Qwen). The sensitivity analysis ($r = 0.80$) mitigates but does not eliminate this concern. Future replications should include baseline conditions for every model tested.

The instructions we plant are ordinary coding preferences, not safety guardrails, and the test turns measure code style rather than a harmful action taken or avoided. We read preferences as a proxy for guardrails because both are user instructions competing for the same attention, but that link is a hypothesis grounded in the shared mechanism, not a demonstrated guardrail bypass. If anything, the proxy may overstate guardrail fragility: real guardrails often sit at position zero in the system prompt, the most stable position (Liu et al.,




Fig. 4. Per-decision posterior predictive check. Observed means (dots) fall within the 90% posterior predictive interval for all 12 decision types.

2024), and are reinforced in training, so it speaks most directly to newly stated, in-context constraints. Showing that a stated safety constraint decays the same way, and that the decay lets a harmful action through, is the experiment this paper motivates but does not run. All of our evidence also comes from Python coding sessions; whether the same per-instruction structure holds for tool-use agents under safety constraints in other domains is untested.

The dataset contains 244 observations, with some decision-by-condition cells as small as 4 observations. The hierarchical model handles sparsity through partial pooling, and the wide credible intervals for low-data decisions reflect this uncertainty. However, the reinforcement ranking for the “uncertain” category (Table 2) should be treated as preliminary.

Compliance was measured at turns 20–25 only. This gives a retention snapshot, not a decay curve. We cannot estimate per-instruction decay rates or say at which turn compliance starts to drop. Fitting forgetting curves would mean scoring every turn, which in turn means designing prompts that let the model demonstrate each decision at each turn without making the conversation artificial.

The deterministic checkers have construct-validity limitations. When a checker cannot find a relevant code pattern, it defaults to a score of 2 (“no clear signal”). Four decision types have default-2 rates above 50%: `architecture_extend` (79%), `docs_regex_comments` (75%), `code_style_listcomp` (72%), and `testing_assert_almost` (50%). The three decisions with the strongest treatment effects (`parametrize`, `standalone`, `no_new`) have 0% default-2 rates, so the headline findings are not driven by this artifact.

This paper does not implement Bayesian Knowledge Tracing. It measures the per-instruction heterogeneity that would make a BKT-based reinforcement system useful. Building and evaluating such a system, including closed-loop token-budget allocation and real-time compliance tracking, is future work.

6.4 Future work

Several directions follow. The most direct is matched baseline and treatment conditions for every model, which would estimate per-model effects directly instead of inferring model-invariance from a single-model check. The checkers, though reproducible, have construct-validity limits, so we would pair them with a blinded panel of cross-family LLM judges and report checker-judge and inter-judge agreement. To separate phrasing from content, we would re-run instructions whose effect excludes zero under negated paraphrases, and the negated one under positive paraphrases. Scoring every turn, rather than only 20–25, would fit per-instruction forgetting curves and decay rates. Finally, running the posterior-derived ranking as a runtime policy would test whether selective reinforcement beats uniform under a fixed token budget, and real coding-session traces would test whether the structure holds outside the harness.

7 Conclusion

We measured per-instruction retention in AI coding assistants and found treatment effects spanning nearly eight log-odds units, from instructions the model adopts only when told, to ones it already follows by default. The pattern appears across the five models and three levels of context pressure tested, and the heterogeneity replicates on a single model with matched baseline and treatment conditions ($r = 0.80$); with baseline data from one model, though, we cannot fully separate per-model from per-instruction effects, and read the model-level result as no detectable effect at this resolution rather than established invariance.

The practical payoff is concrete: of 12 instruction types only 3 show a reinforcement gain whose interval excludes zero in the full model (two when restricted to the single model with matched baseline), so a uniform policy spends most of its tokens where they buy nothing. The retention probabilities estimated here are the input a Bayesian Knowledge Tracing system would need; we leave building it to future work.

Acknowledgments. I thank my partner, Lara Rodrigues La Moglie, for her most gorgeous love and unwavering support; my family, for giving me the opportunities that guided me all the way here; my alma mater, the Universidad Nacional de San Martín, between whose halls I still live; and all my teachers, and the teachers of Argentina, who live each day to inspire others.

Use of large language models During the preparation of this work the author used a large language model (Claude) to improve the manuscript’s language and readability. The author reviewed and edited the text and takes full responsibility for its content.

Disclosure of Interests. The author has no competing interests to declare that are relevant to the content of this article.

References

- Abril-Pla, O., Andreani, V., Carroll, C., Dong, L., Fonnesbeck, C. J., Kochurov, M., Kumar, R., Lao, J., Luhmann, C. C., Martin, O. A., Osthege, M., Vieira, R., Wiecki, T., & Zinkov, R. (2023). PyMC: A modern, and comprehensive probabilistic programming framework in Python. *PeerJ Computer Science*, 9, e1516.
- Chowdhury, B. D. (2026). Lost in the middle at birth: An exact theory of transformer position bias. *arXiv preprint arXiv:2603.10123*.
- Corbett, A. T., & Anderson, J. R. (1994). Knowledge tracing: Modeling the acquisition of procedural knowledge. *User Modeling and User-Adapted Interaction*, 4(4), 253–278.
- Dongre, V., Rossi, R. A., Lai, V. D., Yoon, D. S., Hakkani-Tur, D., & Bui, T. (2025). Drift no more? Context equilibria in multi-turn LLM interactions. *arXiv preprint arXiv:2510.07777*.
- Du, Y., Tian, M., Ronanki, S., Rongali, S., Bodapati, S. B., Galstyan, A., Wells, A., Schwartz, R., Huerta, E. A., & Peng, H. (2025). Context length alone hurts LLM performance despite perfect retrieval. *Findings of the Association for Computational Linguistics: EMNLP 2025*, 23281–23298.
- He, Y., Jin, D., Wang, C., Bi, C., Mandyam, K., Zhang, H., Zhu, C., Li, N., Xu, T., Lv, H., Bhosale, S., Zhu, C., Sankararaman, K. A., Helenowski, E., Kambadur, M., Tayade, A., Ma, H., Fang, H., & Wang, S. (2024). Multi-IF: Benchmarking LLMs on multi-turn and multilingual instructions following. *arXiv preprint arXiv:2410.15553*.
- Hoffman, M. D., & Gelman, A. (2014). The No-U-Turn Sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15, 1593–1623.
- Laban, P., Hayashi, H., Zhou, Y., & Neville, J. (2025). LLMs get lost in multi-turn conversation. *arXiv preprint arXiv:2505.06120*.
- Leviathan, Y., Kalman, M., & Matias, Y. (2025). Prompt repetition improves non-reasoning LLMs. *arXiv preprint arXiv:2512.14982*.
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2024). Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12, 157–173.
- Mu, N., Lu, J., Lavery, M., & Wagner, D. (2025). A closer look at system prompt robustness. *arXiv preprint arXiv:2502.12197*.
- Seyboldt, A. (2024). *Nutpie: A fast sampler for Bayesian posteriors*. <https://github.com/pymc-devs/nutpie>