

Not All Instructions Are Forgotten Equal

Tomás P. Korenblit 

Universidad Nacional de San Martín, Buenos Aires, Argentina
tomaskorenblit@gmail.com

Abstract. Keywords: instruction retention , context rot , Bayesian inference , AI safety , LLM evaluation

No Todas las Instrucciones se Olvidan Igual

Abstract. Keywords: retención de instrucciones , degradación de contexto , inferencia bayesiana , seguridad de la IA , evaluación de LLMs

1 Introduction

2 Related work

3 Method

3.1 Decision planting

3.2 Compliance measurement

3.3 Bayesian model

$$\eta_i = \alpha_{d[i]} + \beta_{d[i]} \cdot \text{treatment}_i$$

Table 1.

Model Size	Baseline	Treatment
------------	----------	-----------

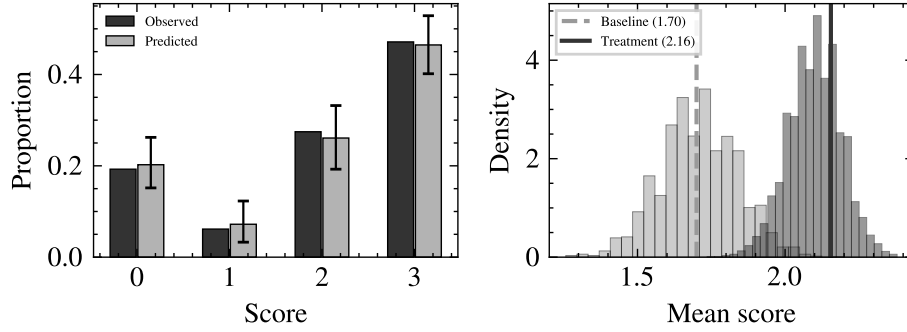


Fig. 1.

4 Experimental setup

4.1 Codebases

4.2 Models

4.3 Dataset

5 Results

5.1 Treatment effect heterogeneity

5.2 Model and codebase effects

5.3 Reinforcement priorities

6 Discussion

6.1 What determines retention

6.2 Robustness checks

6.3 Limitations

6.4 Future work

7 Conclusion

Use of large language models.




Fig. 2.

Table 2.

Decision Gain 94% HDI Action




Fig. 3.




Fig. 4.